



Personality-Based Career Exploration Application

Business Requirements

Created by:	Raymond Palima
Issue date:	June 15, 2026
Version:	1.1

Document Revision History

Release No	Release Date	Author(s)/Editor(s)
1.0	June 1, 2026	Raymond Palima
Summary of Changes/Additions		Initial creation of the document.
1.1	June 4, 2026	Joan Lee Tu, Raymond Palima
Summary of Changes/Additions		Revisions
Summary of Changes/Additions		
Summary of Changes/Additions		
Summary of Changes/Additions		

Table of Contents

1. Key Stakeholders.....	4
2. Project Members.....	4
3. Related Documents.....	4
4. Executive Summary.....	5
5. Project Description.....	6
5.1 Overview/Purpose.....	6
5.2 Future Considerations:.....	6
5.3 Goals and Outcomes.....	6
6. Requirements Determination.....	7
6.1 Functional Requirements (FR):.....	7
6.1.1 Personality Assessment and “Sandbox” FR’s:.....	7
6.1.2 Career Path Search Results FR’s:.....	9
6.1.3 User Access and Academic Pathways FR’s:.....	10
6.2 Non-Functional Requirements:.....	11
6.2.1 Performance and Latency NFR.....	11
6.2.2 Career Path Results NFR’s.....	11
6.2.3 Application Reliability and Capacity.....	12
6.2.4 User interface (UI) and Accessibility.....	12
6.3 Security.....	13
6.3.1 Data Architecture and Implementation Strategy.....	13
6.3.2 AI Agent Data Sanitization:.....	14
6.4 Legal and Compliance.....	14
6.4.1 Minimum Regulatory Compliance (PIPA/PIPEDA).....	14
6.4.2 Algorithmic Liability & AI Disclosure:.....	14
6.4.3 MVP Consent & Deletion.....	15
6.4.4 Terms of Use & Early-Stage Testing Liability Shield.....	15
6.4.5 Required Legal Documentation (To Be Drafted):.....	15
6.4.6 Future Communications (CASL Compliance):.....	15
7. Project Scope.....	16
7.1 In-Scope (Minimum Viable Product - MVP).....	16
7.2 Out-of-Scope.....	17
8. Budget.....	17
9. Proposed Timeline and Deliverables.....	17
9.1 Phase 1: Setup.....	17
9.2 Phase 2: Development of User-Input Form Page.....	18
9.3 Phase 3: Development of the back-end server and AI agent.....	18
9.4 Phase 4: Development of Career Path Results.....	19
9.5 Phase 5: Database, Account Registrations, Security and Legal.....	20
9.6 Phase 6: UI-Revisions, and Internal Bug-testing.....	20
9.7 Phase 7: Deployment and Final Revisions.....	20

9.8 Phase 8: Testing and Updates.....	21
10. Risks and Dependencies.....	21
10.1 Dependencies:.....	21
10.2 Risks:.....	21
11. Assumptions & Constraints.....	22
11.1 Assumptions:.....	22
11.2 Constraints:.....	22
12. Approval.....	22

1. Key Stakeholders

Key Stakeholder	Name
CEO/Project Sponsor	Joan Lee Tu

2. Project Members

Key Stakeholder	Name
Web Developer	Nathan Garcia
Business Analyst/Project Manager	Raymond Palima

3. Related Documents

Document Name	Location
Proof of Concept (Website Demo)	https://career-e-sandbox-poc.vercel.app
Technology Comparison and Selection	 Technology Comparison and Selection
Technical Requirements Document	 Technical Requirements Document
Two-Page Brochure	<i>To be made later.</i>
Slide Deck	<i>To be made later.</i>
How to build your Own - Training Document	<i>To be made later</i>

4. Executive Summary

The Personality-Based Career Exploration Application is a web platform designed to address career and program mismatch, which are commonly experienced by post-secondary students. The platform guides them toward identifying career paths that both align and do not align with their personality or preferences.

The application's core feature is a search engine that processes user inputs from up to nine personality frameworks. This includes models like the Myers-Briggs Type Indicator (MBTI) and the Sparketype Test. The platform is also designed with a "sandbox" philosophy. It allows users to freely adjust specific personality metrics to view alternate career outputs. Results are generated by an Agentic Large Language Model (LLM) and are required to display specific outputs and metrics.

The primary objective is the completion and deployment of a Minimum Viable Product (MVP). Future considerations, although out of scope for the MVP, include integrating prior experience or academic coursework and credit recognition to enhance the search algorithm. Development is constrained by a minimal budget, the non-technical nature of the Project Sponsor, and a very small development team.

5. Project Description

5.1 Overview/Purpose

The purpose of this project is to create a career exploration web application. Unlike other career tests, the platform's search engine uses a combination of personality frameworks (examples include: MBTI and Sparketype) to match users with the most and least fitting career paths. This application will be primarily for education and research purposes. Further development or deployment is contingent upon institutional or sponsored funding.

The primary target market is post-secondary students who are considering a career or program decision. Most of the time, it is because they have not committed to a career path or feel their current path does not fit them. This mismatch often stems from misalignments in personality or preferences. For example, environmental preferences, or social dynamics (e.g., extraversion vs. introversion). This is the problem that the application is addressing. We are providing users with guidance on what career path they should pursue based on their personality.

The platform is designed with a "sandbox" philosophy. For example, the interface should give the user a feeling of creative flexibility. Users can freely toggle, adjust, and modify individual personality metrics without needing to restart from the beginning. This overall design choice encourages curiosity, accommodates users who fall between specific personality archetypes, and allows for free exploration of alternative career matches.

5.2 Future Considerations:

While currently out of scope for the minimum viable product, future versions can incorporate a user's previous academic and professional experience. This can include any courses, diplomas, certificates, etc., already taken. Considering these records allows us to look into transfer equivalences in other programs. This is most useful for enhancing our search algorithm. It also allows users not to be forced into a "hard reset" by being able to transfer their previous classes to a new program. Furthermore, we can even consider the classes they have liked/not liked in our search algorithm. Overall, it would enhance the user experience as it creates more incentive for them to use it. As well as better business opportunities with institute partners. The application can also implement local institutions and their programs as locations users can go to.

5.3 Goals and Outcomes

The ultimate objective of this project is to deploy a fully functional web application. However, we are developing the application in phases. There will be goals in each phase ([refer to timeline below](#)). The primary milestone is the completion of a Minimum Viable Product (MVP).

Upon completion of the MVP, the application will be deployed for external testing. Feedback gathered from the testing will then dictate future updates to the application.

6. Requirements Determination

The following section breaks down the requirements of the application.

6.1 Functional Requirements (FR):

Functional requirements define the specific actions, features, and capabilities the application must perform. “**What** the system should do.”

6.1.1 Personality Assessment and “Sandbox” FR’s:

FR-1. Core Personality Assessment (High Priority):

The system must incorporate 9 personality assessments/models. The overall UI should feel like a single-page form with dropdown navigation for users to input their results into each personality model. From these inputs, the system will determine their desirable/undesirable career paths.

- The supported frameworks are:
 - Myers-Briggs Type Indicator (MBTI - including Assertive/Turbulent)
 - This is the **highest** priority
 - Sparketype Test (Primary and Secondary Sparketypes, and Anti-Sparketype)
 - CliftonStrengths (Top 5 Strengths framework)
 - Big Five Model
 - Enneagram
 - DiSC Assessment
 - Chinese Zodiac, this includes animal and element (Educational/Research purposes)
 - Astrology (Educational/Research purposes)
- Users **cannot** do the personality assessment directly on the platform; they can only enter their results from these assessments.
- The system must incorporate contextual "help" icons or tooltips (e.g., a question mark symbol). These help icons help provide extra details. These can include a short description, alongside external redirection links to official assessment providers. We can also add “good to know” s, like trait implications (e.g., what a high/low score means).

FR-2. “Sandbox” Interface/“Feel” (High Priority):

As mentioned in the project description, the application is supposed to feel like a “sandbox”. The following are meant to address this.

- The system allows the user to “pick and choose” how many personality assessments they want to do (listed above). The user must at least enter the result for one assessment.
- The interface in the career path results must allow users to manually change specific personality traits from their inputs **after** their initial assessment is complete. This allows users to view alternate career outputs without re-taking the assessment (e.g., changing a single score or type).
 - For example, a user is conflicted between being an INTJ or ENTJ (Myers-Briggs Type Indicator). After initially submitting as an INTJ and reviewing their career recommendations, they wish to explore how their results would shift if they were an ENTJ. The platform must accommodate this by allowing the user to adjust that specific trait without needing to restart the assessment process from the beginning.
 - This creates a “sandbox” feel and a creative touch to the platform. It allows users to see changing results for curiosity's sake or if they are unsure of certain personality questions.

FR-3. Additional Optional User Inputs (High Priority):

The following are optional fields that users can input to “enhance” career path results. This section should be a tab that is a long checklist/check box where users can click as many or as few as they want. However, some of these options require you to pick one over the others; thus, the system will limit you to only one. Some inputs are also a slider or a text box.

These options include:

- **Work Environment:** Fully Remote, Hybrid, In-Office.
- **Organizational Structure:** Hierarchical (e.g., Military/Medicine) vs. Flat/Collaborative (e.g., IT).
- **Target Education Level:** A slider ranging from High School Diploma to PhD.
- **Task Dislikes:** A multi-select checklist for user dislikes (e.g., Outdoors, Physical Labour, Heavy Reading/Writing, Science Labs).

- **Demographic Information:** Checklist form for user demographics (e.g., Age, Gender, Race - **If appropriate**)
 - **NOTE:** For this specific option, we should first thoroughly test if it is effective.

For this section, we can utilize AI to cover any fields we are potentially missing. We can also add a textbox where users can add extra details (that we did not cover). However, make sure to thoroughly test this, as it may cause problems.

6.1.2 Career Path Search Results FR's:

FR-4. Core Career Match Results (High Priority)

The system must analyze the personality data inputs provided by the user to generate a customized list of matching career paths. We will utilize an agentic LLM to create an AI agent for this task. We will use one AI model that we determine creates the best results. The following must be incorporated in the results:

- **Percentage Metrics:** The interface must display a specific "compatibility percentage" for each career match based on the user's personality assessment results.
- **Recommended and Non-Recommended List:** The system must display both recommended career paths and non-recommended career paths (e.g. careers to avoid).
- **Reasoning:** Every recommended or non-recommended career path must explain why it does or does not align with the user's personality.
 - This includes the "Pros" and "Cons" for the suggested career path. This is important because while certain career suggestions may hit the majority of the user's requirements, it might still miss out on some.
 - For example, a suggested career meets everything a user wants/needs. However, it is an office job which the user does not like. Based on this, the platform must inform the user of this by indicating it as a "con".

Note: Ideally, we would utilize a completely free LLM. However, we may need to use models that have a cost (determined after testing).

FR-5. Job Market Outlook and Income Level (High Priority)

Every generated career match must include **Job Market Outlook** (growth trends) and **Income/Salary Levels**. Both can be separated into three levels: low, medium and high. For the job market outlook, the levels can be represented as directional arrow signs. While the income level is represented by the “\$” sign (Low: \$, Medium: \$\$, High: \$\$\$). The amount of education required should also be shown.

FR-6. Domain/Industry Filtering (High Priority)

Users must be able to filter their career recommendations by major academic and professional domains (e.g., Medical, Law, Technology, Creative Arts).

FR-7. Storing Career Path Results(Low Priority)

We can store the career path results of users in our database, so they can see them in the future. However, this will take up data space and may cause unwanted costs. This can be part of FR-12.

FR-8. Exporting Career Path Results (Low Priority):

Another option we can use, especially if we cannot store their career path suggestions, is to export a copy of their results, allowing the user to pick among various file formats (e.g. PDF, TXT, CSV, XLS, etc.). This can be an emailed or a printed copy. This is mostly a “quality of life” feature and may take extra time out of the application's core functionality. However, when designing the UI, we can leave the buttons for these features as a “premium” option or to be made later.

FR-9. Switching between Models (Medium Priority):

If possible, we can let users change between LLM models to see different/better results.

FR-10. Backup Models (Low Priority):

Another possibility we can explore is the addition of backup models in case the “main” LLM runs out of tokens.

FR-11. More Details Page (Medium Priority):

An additional feature would be that if a user clicks on a recommended career path, a “more details” page should appear. This page should contain information on the best-fit personality types for that career. It will show a ranking from most compatible to least compatible, and where their personality ranks among them. Furthermore, a dropdown can be implemented to filter it by the top/bottom quartile, or all the personality types from best to worst fit. This gives users full transparency into how their personality stacks compared to everyone else in a career path they are interested in. The most important

personality framework would be the Myers-Briggs and its 16 types, but we could also implement Sparketypes or Astrological signs.

6.1.3 User Access and Academic Pathways FR's:

FR-12. User Registration (High Priority)

The system requires users to register to use the platform. This would mean that we need to set up a database to collect user information when the MVP is released.

The following information must be collected if an account is created:

- Email
- Username

The creation of an account also allows us to add more features (refer to FR-12). It should allow users to test with a limited number of tokens without expending unnecessary tokens.

The application must allow an option for users to opt in for additional “engagement and user feedback” regarding the future progression of the application. Specifically, this includes things like an email list for updates, participating in focus groups/surveys to provide feedback, or booking an appointment for a demo. These choices should be completely optional and are asked during account creation.

FR-13. Optional Paid Tier (Low Priority)

Beyond the MVP, the platform can implement a “paid tier” or a “pay-per-use” for users. This can include multiple features. For example, an optional storage of user personality information. This allows users not to need to re-enter personality assessments every time they use the platform. Exporting career match results could also be a paid feature. Another example would be a more enhanced search that uses a higher-tier AI model.

6.2 Non-Functional Requirements:

Non-functional requirements (NFR) define the system's performance. This includes the quality benchmarks of the platform, its operation, user experience, etc. "**How** the system should do it".

6.2.1 Performance and Latency NFR

NFR-1. AI Response Time (Medium Priority)

Since we are utilizing an Agentic LLM for our results, the application should be able to create results within 60 to 120 seconds (1–2 minutes). However, since we are probably going to utilize free Agentic LLMs, this can be longer (Why it's a medium priority).

6.2.2 Career Path Results NFR's

NFR-2. Loading/Progress Bar (Medium Priority)

In relation to NFR-1, since we are relying on an Agentic LLM for results and thus the results will take a set amount of time, it should show a progress bar to show it is doing something, and the application is not “freezing”. This prevents users from refreshing the page, and allows us to not “run up tokens”.

NFR-3. Recommendation Amount (High Priority)

Every successfully executed search must ideally deliver a minimum of 25 career paths for both recommended and non-recommended. However, if that is not possible, 5-10 would be enough.

NFR-4. Legitimate Data (High Priority)

The results (provided by the LLM) should be legitimate and accurate information. The system must prevent AI hallucinations, ensuring the generated roles, descriptions, and insights are legitimate and accurate.

NFR-5. Results Quality (High Priority)

The results given should be of good quality and applicable. It should be accurate to the user's personality. Exceptions: Chinese Zodiac and Astrology results may be variable.

6.2.3 Application Reliability and Capacity

NFR-6. Application Session Capacity (Low Priority)

The application in the future should be able to handle 100 concurrent active users (depending on the demands of the application). However, for the current development and MVP, it should be able to handle a minimum of 10 users at a time (which should be easily attainable and is why it is a low priority).

NFR-7. Application Availability (High Priority)

Upon the release of the MVP, the system must remain operational with 24/7 availability.

6.2.4 User interface (UI) and Accessibility

NFR-8. UI Clarity (High Priority)

The UI must be responsive and clean. It should be easily understandable and readable. Furthermore, it needs to maintain the “sandbox” feel and functionality. It should also scale cleanly with different desktop resolutions.

Note: Certain future features can be implemented in the UI (Like FR-8), even if they do not work yet. These buttons can say “premium tier”, “to be made later”, etc.

NFR-9. UI Accessibility (Medium Priority)

The interface elements must be accessible (conform to WCAG 2.1 AA accessibility standards). For example, the colour contrast should be sufficient, and the text should not be too small.

6.3 Security

The following security requirements are based on compliance with two major laws: the Personal Information Protection Act (PIPA) and the Personal Information Protection and Electronic Documents Act (PIPEDA). There are more details in [6.4 Legal](#).

6.3.1 Data Architecture and Implementation Strategy

Due to the system requiring mandatory user registration for the platform, a secure database must be established. The security will depend on the chosen infrastructure (see [Technology Comparison and Selection Document](#) for more details)

- **Database options:**

1. **Self-Hosted Relational Database:** This option would involve creating a server on which the database runs. However, if chosen, developers must ensure that server firewalls are configured, SSL certificates are managed, and vulnerabilities are manually patched.
 - **Legal Minimum Requirement:** Under Canadian privacy laws (PIPEDA), storing user passwords in plain text is illegal as it violates the "reasonable safeguards" principle. Therefore, if we manage our own database, we must encrypt/hash passwords. We can easily satisfy this by utilizing standard, free backend libraries (such as bcrypt for Node.js or Python) to hash the passwords in the application code before they are saved to the database.

2. **Managed Database-as-a-Service (DBaaS):** The platform will utilize a third-party database provider to handle baseline security, automatically hashing passwords and providing features like Row Level Security (RLS).
 - *Legal Minimum Requirement:* Utilizing a DBaaS automatically satisfies the PIPEDA "reasonable safeguards" principle regarding passwords right away. Developers must simply utilize the provided RLS features to write strict database rules (e.g., restricting data access so users can only view or modify their own records).
- **Minimum Data Protection (MVP Baseline):** Lastly, the application must use standard HTTPS/SSL encryption to protect data while users are interacting with the website.

6.3.2 AI Agent Data Sanitization:

When transmitting user personality data to the Agentic LLM to create the career path results, the system must ensure that there are **no** personal identifiers (Names, Emails, DOBs, etc.) included. The AI agent must only process raw, anonymous trait markers (e.g., "*User is an MBTI-INFJ and Sparketype-Maven*") to guarantee user privacy.

6.4 Legal and Compliance

This section outlines the mandatory minimum compliance and liability protections required to legally deploy the MVP.

- The two major laws that we are satisfying:
 - [*Personal Information Protection Act \(Alberta\)*](#)
 - [*Personal Information Protection and Electronic Documents Act \(Canada\)*](#)
- However, we also need to keep in mind [*Canada's anti-spam legislation*](#). The current proposed application does not need to comply with this requirement (more information in 6.4.6)

6.4.1 Minimum Regulatory Compliance (PIPA/PIPEDA)

- **Data Minimization:** The system must strictly limit data collection to Username, Email, and personality metrics.
- **Database Location:** To easily comply with Canadian cross-border data laws (PIPEDA) without drafting complex legal transfer agreements, the database server location must be set to Canada.

6.4.2 Algorithmic Liability & AI Disclosure:

The platform must include a clear, visible disclaimer on the results page stating two things:

- Career results and economic metrics are generated by Artificial Intelligence.
- The outputs are for "educational and informational purposes only" and do not constitute certified professional, academic, or financial counselling.

6.4.3 MVP Consent & Deletion

- **Consent Checkbox:** During account creation, users must check a mandatory box agreeing to the Terms of Use and the collection of their Username and Email.
- **Account Deletion:** To legally satisfy a user's right to withdraw consent and have their personal data removed under PIPA/PIPEDA, the application must include a "Delete Account" button within the user settings. This button must permanently wipe the user's data from the database.

6.4.4 Terms of Use & Early-Stage Testing Liability Shield

The application requires a Terms of Use document to be made. Currently, the document must include the following:

- **Beta/Testing Environment Disclaimer:** The platform must be legally designated as an "Early-Stage Research Beta." The Terms of Use will explicitly notify users that the system is operating in a live testing environment to evaluate technical and business model efficiency.
- **Limitation of Liability Clause:** A robust limitation of liability clause must be integrated into the registration/initiation workflow. This clause legally establishes that users access the system, input data, and review AI career recommendations at their own risk, effectively shielding the organization and its research partners from claims arising from data anomalies or software interruptions during the testing and optimization phase.

6.4.5 Required Legal Documentation (To Be Drafted):

While this BRD outlines the necessary technical safeguards, the actual written legal documents (Privacy Policy and Terms of Use) will be drafted at a later stage, before launch. The

development team must include UI placeholders (e.g., footer links) and the required consent checkboxes for these documents so they can be easily linked once finalized.

6.4.6 Future Communications (CASL Compliance):

Currently, the MVP does not include promotional messaging. However, *if* the platform eventually expands to send marketing emails, newsletters, or text messages to registered users, the system must comply with Canada's Anti-Spam Legislation (CASL). This will require adding a separate checkbox specifically for promotional communications (distinct from the Terms of Use checkbox). The platform would also need to ensure that all marketing messages include a functional "unsubscribe" mechanism.

7. Project Scope

Project scope establishes the boundaries of development, dividing requirements into what will be delivered immediately in the Minimum Viable Product (MVP) and what will be delivered after.

7.1 In-Scope (Minimum Viable Product - MVP)

The immediate focus is on creating an MVP. Thus, all the requirements for the MVP are In-Scope (High Priority/"Must have"). The MVP includes:

- **Core Personality Assessment (FR-1 & FR-3):** Integration of personality models and other questions with user inputs.
- **"Sandbox" Interface/"Feel" (FR-2):** The features that make it a sandbox.
- **Core Career Match Results (FR-4):** The deployment of the Agentic LLM framework to generate career paths and match percentages.
- **Job market outlook and filters (FR-5 & FR-6):** Visual indicators for salary ranges, amount of education required, and job market outlook. As well as Domain/Industry filters.
- **User Registration (FR-12):** The development of user registration/accounts and database.
- **Results Quality (NFR-3, NFR-4, NFR-5):** Career path results information quality and integrity.
- **Application Availability (NFR-7):** Application should be available 24/7.
- **Application UI (NFR-8):** Application UI should be understandable, functional, and responsive.
- **Data Privacy and Security (6.3 Security, 6.4 Legal and Compliance):** Security of user data and legal compliance.

The following are still within the scope of the MVP. However, if they are unable to be achieved for the MVP, we can get away with it (Medium Priority/ "Should have"):

- **Switching between Models (FR-9):** If we have multiple LLM models that we deem are exceptional, we can allow users to switch between them to see different/better results.
- **AI Response Time (NFR-1):** Result response time from AI within 60 to 120 seconds (1–2 minutes). Note: this is reliant on the AI model we are using, and since we are utilizing free models, it may take longer (Reason as to why it is medium priority/“should have”).
- **Results Progress Bar (NFR-2):** Progress bar for the career path results response.
- **UI Accessibility (NFR-9):** Application must be accessible.
- **More Details Page (FR-11):** An additional page where users can click on a specific career path to see how their personality ranks compared to everyone else.

7.2 Out-of-Scope

Features not needed for the MVP, and can be implemented in the future (Low Priority/ “Nice to Have”):

- **Paid Tier (FR-13):** An optional tier that users can pay for.
- **Results (FR-7 & FR-8):** Storage of the user’s career path results, as well as exporting files or sending a copy to their email.
- **High-Volume Server Scalability (NFR-6):** Architecture expansion to handle 100+ concurrent users (if there is a demand for it, as the MVP will probably not need it).
- **Backup Models (FR-10):** The addition of backup LLM models that users can use in case the main models run out of tokens.

8. Budget

The budget for this project would be to keep the cost at a minimum, thus meaning to use as many cost-free options as possible. Regarding the future, the project is hoping to get funding/support from post-secondary institutions. If we are able to gain funding, we can then utilize it to develop the project.

An Important note is that a significant part of our budget relies on the AI we will use. Ideally, we will utilize a completely free LLM. However, we may need to use models that have a cost (determined after testing).

9. Proposed Timeline and Deliverables

9.1 Phase 1: Setup

Timeline: 0.5 - 1 Week

Deliverables and Focus:

- Decide on the technical stack (what language to use, database, server, etc.)

- While not **all** of it has to be decided, the front-end language/technology needs to be decided.
 - Refer to the [Technology Comparison and Selection](#) document and the [Technical Requirements Document](#).
- Initialize the setup
 - The main focus is just setting up the front-end on the local computers of the developers. However, you can set up everything else if we have decided on the tech stack (Back-end server, database, etc.).
 - Reminder to create a repository for the project (GitHub).
- Initialize the front-end
 - Set up the basic user interface (UI) of the overall look and feel of the webpage (Homepage, navigation bar, colour theme of website, etc.).

9.2 Phase 2: Development of User-Input Form Page

Timeline: 1 - 2+ Weeks

Deliverables and Focus:

- Determine the UI of the “form.”
 - Design and build the form (dropdown style) that the user will enter their personal information.
- Decide on what information should be asked from the user regarding the personality assessments/models (Myers-Briggs, Sparketype, etc.).
 - Research into the results of these personality frameworks and determine what information we need to ask the user (what should we ask? how should we ask? etc.)
- Format the form based on our research.
 - Once we determine what we need to ask users, we need to design and add the appropriate questions/fields to the form UI.
- Incorporate the Additional Optional User Inputs (FR-3).

9.3 Phase 3: Development of the back-end server and AI agent.

Timeline: 1 - 2+ Weeks

Deliverables and Focus:

- Set up the back-end server if we have not set it up yet.
 - Note: At this point, we are still just running the servers on local machines, like the front-end.
- Incorporate the AI agent.
 - Connect an AI agent (whichever we decide to test first) to our backend server.
 - Set up/generate all the API keys needed.
 - Test to see if we can get a response from the AI to test the connection.
- “Train” the AI agent.
 - Give the AI agent context (its purpose, what information it’ll take, what the export should be, etc.)
 - At this point, we can just let the AI show unstructured text information (Not formatted properly for our UI) to see its response.
 - Test and improve the AI response.
 - One way to improve it is by altering the context we provide it.
- Test other AI LLM’s to find the best result.
 - Test other LLM’s to see which is the best to utilize for the platform (basically repeating steps 2 and 3).
- Decide on which LLM to use for the platform.

9.4 Phase 4: Development of Career Path Results.

Timeline: 0.5 - 1 Week

Deliverables and Focus:

- Design and create the UI for the career path results.
- Connect the result of the AI agent to our front-end UI.
 - Simplest terms: fit the AI results into the structure and format of our website.

- Ensure that the results are meeting all our requirements.

9.5 Phase 5: Database, Account Registrations, Security and Legal.

Timeline: 1 Week

Deliverables and Focus:

- Decide on final tech structure if we haven't (creating our own server, or utilizing 3rd party servers/database, etc.)
- Setup database if we have not done so already.
- Incorporate the creation of user accounts and profiles.
- User Authentication.
- Add "search" limits per account.
- Allow the deletion and modification of user accounts.
- Satisfy/implement the security and legal needs to satisfy our requirements.
 - This includes adding the terms of use and privacy policy.
 - Reminder to ensure that no user-sensitive information is being sent to the AI-agent.

9.6 Phase 6: UI-Revisions, and Internal Bug-testing.

Timeline: 0.5 - 1 Week

Deliverables and Focus:

- Polish up UI.
- Finish all requirements.
- Internal bug testing.

9.7 Phase 7: Deployment and Final Revisions.

Timeline: 0.5 Week

Deliverables and Focus:

- Deployment to the public.
- Final bug tests (new ones typically come out after deployment).
- Final Revisions before public testing.

9.8 Phase 8: Testing and Updates

Timeline: No-timeline

Deliverables and Focus:

- Gain feedback from users.
- Update the application based on feedback from users.

10. Risks and Dependencies

This section highlights what can go wrong and what factors the MVP relies on.

10.1 Dependencies:

- **Third-Party AI Lock-in:** The career path search is entirely reliant on an external Agentic LLM API (E.g., OpenAI, Anthropic). If the provider experiences an outage or changes the availability, the platform will be unavailable.
- **External Personality Assessments:** The platform assumes and is reliant on users independently taking their personality tests (MBTI, Sparketype, etc.) via external, third-party websites before manually inputting their results into the system.

10.2 Risks:

- **AI Hallucinations, Response and Formatting Failures:** The LLM may occasionally return unformatted text, hallucinated/incorrect data. This could provide illogical career advice or break the UI of the platform.
- **API Rate Limits and Token Cost:** Relying on free API tiers for the LLM or Managed Database (DBaaS) may result in the system crashing or blocking users if we exceed concurrent user limits. Furthermore, if the free tiers are ineffective, we may also require a higher tier that has a cost.

- **Timeline delay:** Phases 2 and 3 of the proposed timeline may take longer than expected, as they are the heaviest and longest phases of the application.

11. Assumptions & Constraints

This section defines the boundaries of the platform.

11.1 Assumptions:

- **User Honesty:** The system assumes users will accurately and honestly input their true personality test results. We are not verifying their scores against the external testing platforms.
- **LLM:** The system assumes that an LLM will be able to satisfy the requirements for the career path results.

11.2 Constraints:

- **Strict Budget/Resources:** As mentioned before, the platform is constrained by a minimal budget. This means the team will use as many cost-free resources as possible.
- **Project Sponsor:** The project sponsor is not technical.
- **Team Size:** The development team for the application is very small.

12. Approval

Authority	Name	Role	Date
Not started ▾	Joan Lee Tu	CEO/Project Sponsor	